




Bare encrypted Docker endpoint setup: concurrency sweep

Workers = endpoint setups running concurrently. Subprocess slots = external commands running concurrently.
Native Netlink calls do not use subprocess slots. Example: 32/8 allows 32 setups, with at most 8 commands at once.




256 endpoints per cell; cached Docker PID; iptables outside timer; repeated sysctl retained.
Each endpoint has its own VXLAN, bridge, veth pair, UDP port and four XFRM objects.
Blank cells were not tested. Short runs screen candidates; use 1,000-endpoint repeats for selection.

PID/iptables-only changes: 8/8 is a practical candidate




Bars: median; whiskers: observed min/max, not confidence intervals. Slots above workers are grouped.
All runs check object counts, Docker PID/start-time continuity, and removal of test interfaces.
Rates are endpoint creation throughput, not complete cross-host connections or HTTP requests.

Required operations, fewer subprocesses



Two 1,000-endpoint runs per case and host. Namespace batch preserves address-before-UP; XFRM batch preserves the four state/policy operations and order. No Nullnet code changes. Bare setup checks do not replace encrypted traffic, restart, partial-failure and lifecycle testing.



Two 1,000-endpoint trials per pair and host, ordered 8 / 32 / 32 / 8 workers. Both use 8 command slots. Whiskers show observed min/max. The lab is KVM-based and later sweep trials slowed on both nodes; 32/8 was more consistent here; 8/8 matched it once. Neither is proven a universal optimum.